

# K-Means Clustering

Enhancing performance and scalability through parallelism using  
MPI

Luís Neto (a77763), Gonçalo Raposo (a77211)

January 4, 2019

## 1 Case of Study

*K-means clustering* is an algorithm commonly used to partition data into  $k$  groups based on their similarities. The algorithm consists in three step:

1. Selection of  $k$  random centroids.
2. Assignment Step: each point is assigned to a cluster based on closest centroid determined with *Euclidean distance*.
3. Update Step: each centroid is updated to be the mean of his cluster.

The latter two phases are repeated until the centroids no longer change.

## 2 Distributed Memory Implementation

### 2.1 Partition

In this development phase, we started by decomposing the algorithm in different phases taking into account the functional decomposition and the data domain of what could be parallelized.

The program begins by finding the maximum value of the data set. Based on the maximum value, we generate  $k$  random centroids, being so, the data domain is the data set points.

Afterward, we calculate the minimum distance between each point and centroids. Subsequently, all the points are assigned to the correspondent cluster. In this phase, it is fairly easy to understand that the data domain are the points from the data set.

In the third phase of the algorithm, we calculate the error and update all centroids. Since each centroid is the average of his cluster, the data domain is the data set's points.

### 2.2 Communication

As said before, this algorithm has 3 phases that need to execute in order, due to the data dependencies. After the first phase, the process with rank 0 will communicate the data set points and centroids to the workers.

In the end of phase 2, each process shares his local clusters and updates them, with the information received from the other processes.

In the last phase, we calculate part of the error and update the centroids, with the local information, afterwards each process shares his centroids and updates them, with the information received from the other processes.

Finally, we calculate the final error and send a final message from the process with rank 0 to the other processes. This message will continue or stop the algorithm. When the algorithm converges all Workers send a message to the process with rank 0 containing the current sets, otherwise, we restart the iterative process.

### 2.3 Agglomeration

In this phase, we consider that the granularity is a chunk of the data set with size equal to  $\frac{datasetSize}{\#Processes}$ . With this granularity, we can use several collective operations to increase performance and scalability.

In the first communication, we choose to use *Scatter* and *Broadcast* to divide the data set by the processes and to send the centroids, respectively.

Between phase 2 and 3, we tested 2 strategies, a *Reduce* followed by a *Broadcast* and a *Allreduce*, this strategies are used to share and update the local size counters of each cluster. The efficiency of these operations depends directly from the OpenMPI implementation, however, the latter strategy should be better than the former. In the next sections we will refer to this communication as comm2.

In order to calculate the final error, we need to send the information of each local updated centroid, therefore, we also tested the 2 strategies described above. In the next sections we will refer to this communication as comm3.

Finally, to stop the algorithm we *Broadcast* a message and if this message is positive we use an *Scatter* operation to get the clusters of process, in contrast, if the message is negative, we continue the algorithm. In the next section we will refer to this communications as error\_check.

### 2.4 Task Mapping

As explained before, in phase 2 we assign each point to a cluster based on closest centroid and in Phase 3 each centroid is updated to the average of his cluster, so, the workload for each element of data, centroid or point, is the same. With this we can do a regular partition of the data set and pass it to each process.

Since the algorithm is iterative, being the local centroids updated with the last iteration cluster's points, we can say that this algorithm fits in the category of the Heartbeat Algorithms.

### 2.5 Algorithm Analysis

In this section, we describe the sequential and parallel algorithm's complexity, N will represent the data set size, K the number of clusters and P the number of processes.

For the sequential version, in the first phase we iterate the data set in order to obtain the maximum value, subsequently, we generate all centroids, therefore, this phase complexity is  $N + K$ .

Phase 2 calculates the minimum distance between each point and centroids, consequentially, this phase complexity is  $N * K$ .

Finally, for phase 3, we iterate the clusters in order to calculate the first part of the error, update the clusters by iterating the data set and caculate the final error by iterating the clusters, hence, this phase complexity is  $2 * K + N$ . With this we can say that the sequential algorithm complexity is:

$$(N + K) + (N * (K + 1) + 2 * K) \quad (1)$$

In the parallel version, we equally divide the data set by the processes, so, we can rewrite the last complexity as:

$$(N + K) + (\frac{N}{P} * (K + 1) + 2 * K) \quad (2)$$

In this version, we also need to take in account the communication complexity, we will consider the following values for the primitives used.

1. Broadcast :  $\log_2(P)$
2. Reduce :  $\log_2(P)$
3. Scatter :  $P$
4. Gather :  $P$

After the first phase, we send the data set chunk to each process and clusters using the primitive Scatter and Broadcast, hence, the communication complexity is  $2 * \log_2(P) + P$ .

At the end of second phase, we reduce the local counters of each cluster to one process and broadcast the result, therefore, the complexity is  $2 * \log_2(P)$ .

After the update the local clusters, we reduce the local clusters to one process and broadcast the result, therefore, the complexity is  $4 * \log_2(P)$ .

Finally, after the final error calculation, we broadcast a message to each process in order to stop or continue the algorithm. If the algorithm continues the complexity is  $\log_2(P)$ , otherwise, the complexity is  $P + \log_2(P)$ , due to the fact that we use a *Gather* primitive to collect the final result.

Having the communication complexity we can write the parallel algorithm complexity as:

$$(N + K) + (\frac{N}{P} * (K + 1) + 2 * K) + 9 * \log_2 P + 2 * P \quad (3)$$

With this the speedup can be obtained by:

$$\frac{(N + K) + (N * (K + 1) + 2 * K)}{(N + K) + (\frac{N}{P} * (K + 1) + 2 * K) + 9 * \log_2(P) + 2 * P} \quad (4)$$

### 3 Input tests description

In order to test the developed algorithms, 2 different data sets were created by a Python script using a real uniform distribution, the first data set has 1966080 points and the last one 62914560. We choose the size of the first data set, in order to, fully occupy the level 3 cache. The other data set size was chosen, so that, it won't fit in any cache level.

## 4 Results Analysis

### 4.1 Communication

In order to choose the best mapping and communication strategy, we built 3 figures where we can observe the overall communication time, the partial communication time and the time spent on CPU and communications.

#### 4.1.1 Overall Communication Time

By observing the figure, is possible to see that, in terms of communication, the mapping by core is better up to 16 core, however, when we use 2 nodes, the time spent of this mapping increases and it is similar when we use 16 cores in both machines. It also possible to identify that *comm1* is the most expensive communication.

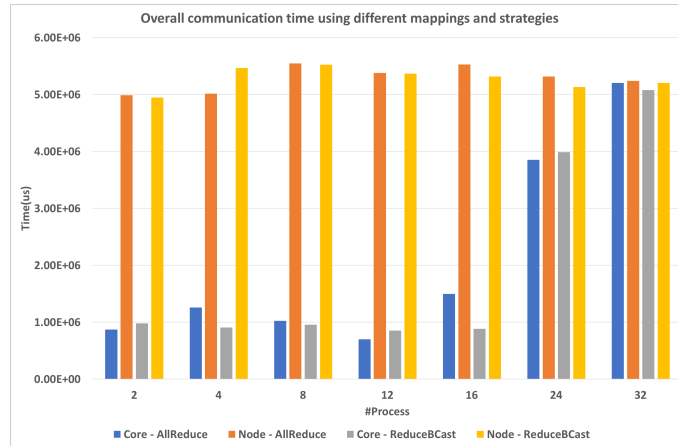


Figure 1: COMM time(us) with different mappings and strategies for the larger data set

#### 4.1.2 Partial Communication Time

In this figure, only the time spent in *comm2* and *comm3* are presented, so that we can compare both strategies and mappings. Between 2 and 32 processes, in both mappings, it is possible to see that for some #processes it is better to use *ReduceBroadcast* and for others to use *Allreduce*, however, the former strategy, in average, is better than the latter.

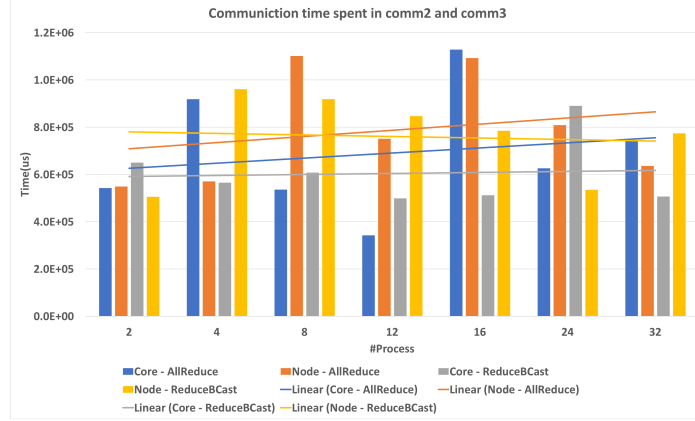


Figure 2: Communication time in *comm2* and *comm3*

## 4.2 Execution Profile

With this section, we will describe the algorithm's execution profile. As referred in section 2, this algorithm executes 4 communications and has 3 computation phases. The time spent in CPU is expected to be greater than the communication time, due to the fact that between iterations we only pass small messages. These messages either are the local sets counters or the centroids, with  $clusters * sizeof(double)$  and  $2 * clusters * sizeof(double)$  bytes, respectively.

Figure 2 represents the median time spent in each phase and in each communication. In order to obtain these results, we used the larger data set, core mapping and *Reduce + Broadcast* strategy.

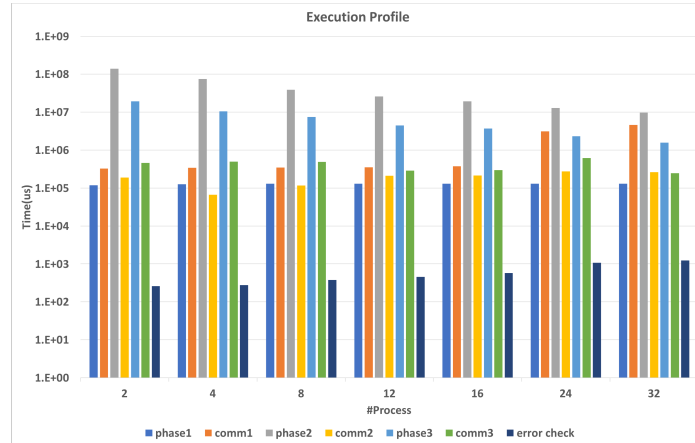


Figure 3: Execution profile for using the larger data set, core mapping and *Reduce + Broadcast* strategy.

As expected, when we use up to 16 processes, phase 2 and 3 times are greater than the communication times. Beyond that number communications will start to increase, since we start to use 2 computing platforms.




Figure 4: Execution profile for using the larger data set, core mapping and *Reduce + Broadcast* strategy.

### 4.3 Load Balance

In section D of the appendix, there are all the tables with the execution time of each process, for each possible combination of mapping, strategy, and size.

From the data in these tables, we can confirm that our algorithm has a good load balance, being the time of execution in each process approximately equal to  $\frac{Totaltime}{\#process}$ .

### 4.4 Achieved Performance

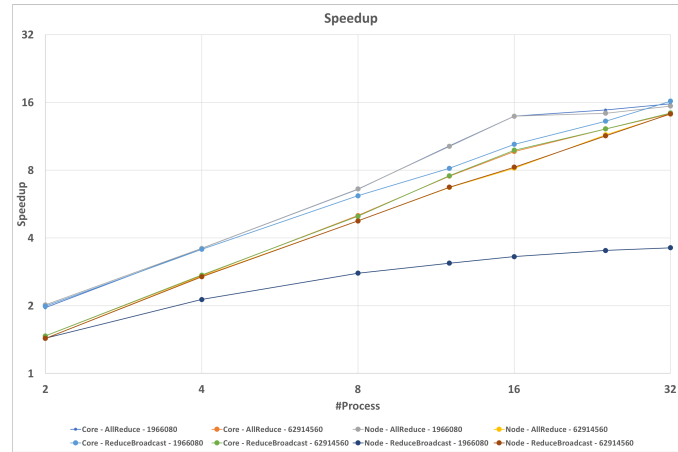


Figure 5: Speedup of each possible combination of size, mapping and strategy.

#### 4.4.1 Data set with 1966080 points

For this data set, it is possible to identify that the worst mapping is to map the processes by node and use the *ReduceBroadcast* method. Meanwhile, all the other strategies implemented have similar behaviors scaling well up to 16 processes and speedups close to the theoretical ones.

#### 4.4.2 Data set with 62914560 points

In respect to this data set, its general behavior is equivalent of the data set previously analyzed. In this case, the strategy of using *ReduceBroadcast* while applying the processes by each node available is the worst among all techniques. All the other, have similar speedups gains by the number of processes. To note that the speedup obtained is almost half by every process added to the task.

## 5 Conclusion

With the performance results obtained, we concluded that the best strategy to implement the *k-means* algorithm using a distributed memory approach is to map all the processes by core available while using the *ReduceBroadcast* collective primitive. We can also conclude that both paradigms have similar speedup, however, if the data set is larger than the RAM of one machine, only the MPI version will allow us to solve the problem.

Finally, taking in account the cost of implementation, the distributed memory paradigm is easier to program and the code is more readable.

# Appendix

## A Computing Platforms Hardware characterization

|                     |                            |
|---------------------|----------------------------|
| Manufacturer        | Intel Corporation          |
| Architecture        | Ivy Bridge                 |
| Model               | Xeon E5-2650v2             |
| # Cores             | 16                         |
| # Threads           | 32                         |
| Processor Frequency | 2.6 GHz                    |
| L1 Cache            | 32 KB (Data + Instruction) |
| L2 Cache            | 256 KB                     |
| L3 Cache            | 20 MB                      |
| RAM Memory          | 64Gb                       |
| Network             | gbe/myri                   |

Table 1: Nodes 641 hardware characterization

## B Node Mappings

### B.1 By Node

### B.1.1 2 Processes

```
[compute-641-8.local:06519] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../../][../../../../../]
[compute-641-20.local:16484] MCW rank 1 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../../][../../../../../]
```

### B.1.2 4 Processes

```
[compute-641-8.local:06526] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../.././../.././../.././][../.././../.././../.././]
[compute-641-8.local:06526] MCW rank 2 bound to socket 0[core 1[hwt 0-1]]:
[../../BB/../.././../.././../.././][../.././../.././../.././../.././]
[compute-641-20.local:16490] MCW rank 1 bound to socket 0[core 0[hwt 0-1]]:
[BB/../.././../.././../.././../.././][../.././../.././../.././../.././]
[compute-641-20.local:16490] MCW rank 3 bound to socket 0[core 1[hwt 0-1]]:
[../../BB/../.././../.././../.././][../.././../.././../.././../.././]
```

### B.1.3 8 Processes

```
[compute-641-8.local:06532] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
```



```

[compute-641-8.local:06532] MCW rank 2 bound to socket 0[core 1[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-8.local:06532] MCW rank 4 bound to socket 0[core 2[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-8.local:06532] MCW rank 6 bound to socket 0[core 3[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16496] MCW rank 1 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../../../../../]
[compute-641-20.local:16496] MCW rank 3 bound to socket 0[core 1[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16496] MCW rank 5 bound to socket 0[core 2[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16496] MCW rank 7 bound to socket 0[core 3[hwt 0-1]]:
[..../BB/../../../../../../../../]

```

#### B.1.4 12 Processes

```

[compute-641-8.local:06542] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../../../../../]
[compute-641-8.local:06542] MCW rank 2 bound to socket 0[core 1[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-8.local:06542] MCW rank 4 bound to socket 0[core 2[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-8.local:06542] MCW rank 6 bound to socket 0[core 3[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-8.local:06542] MCW rank 8 bound to socket 0[core 4[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-8.local:06542] MCW rank 10 bound to socket 0[core 5[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16506] MCW rank 1 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../../../../../]
[compute-641-20.local:16506] MCW rank 3 bound to socket 0[core 1[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16506] MCW rank 5 bound to socket 0[core 2[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16506] MCW rank 7 bound to socket 0[core 3[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16506] MCW rank 9 bound to socket 0[core 4[hwt 0-1]]:
[..../BB/../../../../../../../../]
[compute-641-20.local:16506] MCW rank 11 bound to socket 0[core 5[hwt 0-1]]:
[..../BB/../../../../../../../../]

```

#### B.1.5 16 Processes

```

[compute-641-8.local:06556] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../../../../../]
[compute-641-8.local:06556] MCW rank 2 bound to socket 0[core 1[hwt 0-1]]:
[..../BB/../../../../../../../../]

```

```

[compute-641-8.local:06556] MCW rank 4 bound to socket 0[core 2[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06556] MCW rank 6 bound to socket 0[core 3[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06556] MCW rank 8 bound to socket 0[core 4[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06556] MCW rank 10 bound to socket 0[core 5[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06556] MCW rank 12 bound to socket 0[core 6[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06556] MCW rank 14 bound to socket 0[core 7[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 1 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 3 bound to socket 0[core 1[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 5 bound to socket 0[core 2[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 7 bound to socket 0[core 3[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 9 bound to socket 0[core 4[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 11 bound to socket 0[core 5[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 13 bound to socket 0[core 6[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-20.local:16519] MCW rank 15 bound to socket 0[core 7[hwt 0-1]]:
[../BB/../../..][../..]

```

#### B.1.6 24 Processes

```

[compute-641-8.local:06575] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 2 bound to socket 0[core 1[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 4 bound to socket 0[core 2[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 6 bound to socket 0[core 3[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 8 bound to socket 0[core 4[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 10 bound to socket 0[core 5[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 12 bound to socket 0[core 6[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 14 bound to socket 0[core 7[hwt 0-1]]:
[../BB/../../..][../..]
[compute-641-8.local:06575] MCW rank 16 bound to socket 1[core 8[hwt 0-1]]:
[../BB/../../..][BB/../../..]

```



```

[compute-641-8.local:06601] MCW rank 16 bound to socket 1[core 8[hwt 0-1]]:
[../..../..../..../..][BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 18 bound to socket 1[core 9[hwt 0-1]]:
[../..../..../..../..][../BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 20 bound to socket 1[core 10[hwt 0-1]]:
[../..../..../..../..][../..BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 22 bound to socket 1[core 11[hwt 0-1]]:
[../..../..../..../..][../..../BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 24 bound to socket 1[core 12[hwt 0-1]]:
[../..../..../..../..][../..../..BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 26 bound to socket 1[core 13[hwt 0-1]]:
[../..../..../..../..][../..../..../BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 28 bound to socket 1[core 14[hwt 0-1]]:
[../..../..../..../..][../..../..../..BB/..../..../..../..]
[compute-641-8.local:06601] MCW rank 30 bound to socket 1[core 15[hwt 0-1]]:
[../..../..../..../..][../..../..../..../BB/..../..../..../..]
[compute-641-20.local:16562] MCW rank 1 bound to socket 0[core 0[hwt 0-1]]:
[BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 3 bound to socket 0[core 1[hwt 0-1]]:
[../BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 5 bound to socket 0[core 2[hwt 0-1]]:
[../..BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 7 bound to socket 0[core 3[hwt 0-1]]:
[../..../BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 9 bound to socket 0[core 4[hwt 0-1]]:
[../..../..BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 11 bound to socket 0[core 5[hwt 0-1]]:
[../..../..../BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 13 bound to socket 0[core 6[hwt 0-1]]:
[../..../..../..../BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 15 bound to socket 0[core 7[hwt 0-1]]:
[../..../..../..../..BB/..../..../..../..][../..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 17 bound to socket 1[core 8[hwt 0-1]]:
[../..../..../..../..][BB/..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 19 bound to socket 1[core 9[hwt 0-1]]:
[../..../..../..../..][../BB/..../..../..../..../..]
[compute-641-20.local:16562] MCW rank 21 bound to socket 1[core 10[hwt 0-1]]:
[../..../..../..../..][../..BB/..../..../..../..]
[compute-641-20.local:16562] MCW rank 23 bound to socket 1[core 11[hwt 0-1]]:
[../..../..../..../..][../..../BB/..../..../..../..]
[compute-641-20.local:16562] MCW rank 25 bound to socket 1[core 12[hwt 0-1]]:
[../..../..../..../..][../..../..../BB/..../..../..../..]
[compute-641-20.local:16562] MCW rank 27 bound to socket 1[core 13[hwt 0-1]]:
[../..../..../..../..][../..../..../..BB/..../..../..../..]
[compute-641-20.local:16562] MCW rank 29 bound to socket 1[core 14[hwt 0-1]]:
[../..../..../..../..][../..../..../..../BB/..../..../..../..]
[compute-641-20.local:16562] MCW rank 31 bound to socket 1[core 15[hwt 0-1]]:
[../..../..../..../..][../..../..../..../..BB/..../..../..../..]

```

## B.2 By Core

### B.2.1 2 Processes

```
[compute-641-8.local:04530] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
```

```
[BB/./././././././././.[./././././././././.]
```

```
[compute-641-8.local:04530] MCW rank 1 bound to socket 0[core 1[hwt 0-1]]:
```

```
[./BB/././././././././.[./././././././././.]
```

### B.2.2 4 Processes

```
[compute-641-8.local:04536] MCW rank 0 bound to socket 0[core 0[hwt0-1]]:
[BB/../../../../][../../../../]
[compute-641-8.local:04536] MCW rank 1 bound to socket 0[core 1[hwt 0-1]]:
[./BB/../../../../][../../../../]
[compute-641-8.local:04536] MCW rank 2 bound to socket 0[core 2[hwt 0-1]]:
[../../../../BB/../../../../][../../../../]
[compute-641-8.local:04536] MCW rank 3 bound to socket 0[core 3[hwt 0-1]]:
[../../../../BB/../../../../][../../../../]
```

### B.2.3 8 Processes

```
[compute-641-8.local:04546] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 1 bound to socket 0[core 1[hwt 0-1]]:
[./BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 2 bound to socket 0[core 2[hwt 0-1]]:
[././BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 3 bound to socket 0[core 3[hwt 0-1]]:
[./././BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 4 bound to socket 0[core 4[hwt 0-1]]:
[././././BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 5 bound to socket 0[core 5[hwt 0-1]]:
[./././././BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 6 bound to socket 0[core 6[hwt 0-1]]:
[././././././BB/../../../../][./../../../../]
[compute-641-8.local:04546] MCW rank 7 bound to socket 0[core 7[hwt 0-1]]:
[./././././././BB/../../../../][./../../../../]
```

### B.2.4 12 Processes

```
[compute-641-8.local:04564] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/../../../../..][../../../../..]
[compute-641-8.local:04564] MCW rank 1 bound to socket 0[core 1[hwt 0-1]]:
[../../../../BB/../../../../..][../../../../..]
[compute-641-8.local:04564] MCW rank 2 bound to socket 0[core 2[hwt 0-1]]:
[../../../../BB/../../../../..][../../../../..]
[compute-641-8.local:04564] MCW rank 3 bound to socket 0[core 3[hwt 0-1]]:
```

```

[../././BB/././././][.././././././././]
[compute-641-8.local:04564] MCW rank 4 bound to socket 0[core 4[hwt 0-1]]:
[.././././BB/./././][.././././././././]
[compute-641-8.local:04564] MCW rank 5 bound to socket 0[core 5[hwt 0-1]]:
[../././././BB/././][.././././././././]
[compute-641-8.local:04564] MCW rank 6 bound to socket 0[core 6[hwt 0-1]]:
[.././././././BB/./][.././././././././]
[compute-641-8.local:04564] MCW rank 7 bound to socket 0[core 7[hwt 0-1]]:
[../././././././BB][.././././././././]
[compute-641-8.local:04564] MCW rank 8 bound to socket 1[core 8[hwt 0-1]]:
[.././././././././][BB/./././././././]
[compute-641-8.local:04564] MCW rank 9 bound to socket 1[core 9[hwt 0-1]]:
[.././././././././][../BB/././././././]
[compute-641-8.local:04564] MCW rank 10 bound to socket 1[core 10[hwt 0-1]]:
[.././././././././][.././BB/./././././]
[compute-641-8.local:04564] MCW rank 11 bound to socket 1[core 11[hwt 0-1]]:
[.././././././././][../././BB/././././]

```

### B.2.5 16 Processes

```

[compute-641-8.local:04590] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/././././././][.././././././././]
[compute-641-8.local:04590] MCW rank 1 bound to socket 0[core 1[hwt 0-1]]:
[../BB/./././././][.././././././././]
[compute-641-8.local:04590] MCW rank 2 bound to socket 0[core 2[hwt 0-1]]:
[.././BB/././././][.././././././././]
[compute-641-8.local:04590] MCW rank 3 bound to socket 0[core 3[hwt 0-1]]:
[../././BB/./././][.././././././././]
[compute-641-8.local:04590] MCW rank 4 bound to socket 0[core 4[hwt 0-1]]:
[.././././BB/././][.././././././././]
[compute-641-8.local:04590] MCW rank 5 bound to socket 0[core 5[hwt 0-1]]:
[../././././BB/./][.././././././././]
[compute-641-8.local:04590] MCW rank 6 bound to socket 0[core 6[hwt 0-1]]:
[.././././././BB/./][.././././././././]
[compute-641-8.local:04590] MCW rank 7 bound to socket 0[core 7[hwt 0-1]]:
[../././././././BB][.././././././././]
[compute-641-8.local:04590] MCW rank 8 bound to socket 1[core 8[hwt 0-1]]:
[../././././././][BB/././././././]
[compute-641-8.local:04590] MCW rank 9 bound to socket 1[core 9[hwt 0-1]]:
[../././././././][../BB/./././././]
[compute-641-8.local:04590] MCW rank 10 bound to socket 1[core 10[hwt 0-1]]:
[../././././././][.././BB/././././]
[compute-641-8.local:04590] MCW rank 11 bound to socket 1[core 11[hwt 0-1]]:
[../././././././][../././BB/./././]
[compute-641-8.local:04590] MCW rank 12 bound to socket 1[core 12[hwt 0-1]]:
[../././././././][.././././BB/././]
[compute-641-8.local:04590] MCW rank 13 bound to socket 1[core 13[hwt 0-1]]:
[../././././././][../././././BB/./]
[compute-641-8.local:04590] MCW rank 14 bound to socket 1[core 14[hwt 0-1]]:

```



```

[../..../BB/..][../..../..../..]
[compute-641-20.local:15351] MCW rank 22 bound to socket 0[core 6[hwt 0-1]]:
[../..../BB/..][../..../..../..]
[compute-641-20.local:15351] MCW rank 23 bound to socket 0[core 7[hwt 0-1]]:
[../..../BB/..][../..../..../..]

```

### B.2.7 32 Processes

```

[compute-641-8.local:04660] MCW rank 0 bound to socket 0[core 0[hwt 0-1]]:
[BB/..../..../..][../..../..../..]
[compute-641-8.local:04660] MCW rank 1 bound to socket 0[core 1[hwt 0-1]]:
[../BB/..../..../..][../..../..../..]
[compute-641-8.local:04660] MCW rank 2 bound to socket 0[core 2[hwt 0-1]]:
[../BB/..../..../..][../..../..../..]
[compute-641-8.local:04660] MCW rank 3 bound to socket 0[core 3[hwt 0-1]]:
[../..../BB/..../..][../..../..../..]
[compute-641-8.local:04660] MCW rank 4 bound to socket 0[core 4[hwt 0-1]]:
[../..../BB/..../..][../..../..../..]
[compute-641-8.local:04660] MCW rank 5 bound to socket 0[core 5[hwt 0-1]]:
[../..../BB/..../..][../..../..../..]
[compute-641-8.local:04660] MCW rank 6 bound to socket 0[core 6[hwt 0-1]]:
[../..../BB/..][../..../..../..]
[compute-641-8.local:04660] MCW rank 7 bound to socket 0[core 7[hwt 0-1]]:
[../..../BB/..][../..../..../..]
[compute-641-8.local:04660] MCW rank 8 bound to socket 1[core 8[hwt 0-1]]:
[../..../BB/..../..][BB/..../..../..]
[compute-641-8.local:04660] MCW rank 9 bound to socket 1[core 9[hwt 0-1]]:
[../..../BB/..../..][../BB/..../..../..]
[compute-641-8.local:04660] MCW rank 10 bound to socket 1[core 10[hwt 0-1]]:
[../..../BB/..../..][../BB/..../..../..]
[compute-641-8.local:04660] MCW rank 11 bound to socket 1[core 11[hwt 0-1]]:
[../..../BB/..../..][../BB/..../..../..]
[compute-641-8.local:04660] MCW rank 12 bound to socket 1[core 12[hwt 0-1]]:
[../..../BB/..../..][../BB/..../..../..]
[compute-641-8.local:04660] MCW rank 13 bound to socket 1[core 13[hwt 0-1]]:
[../..../BB/..../..][../BB/..../..../..]
[compute-641-8.local:04660] MCW rank 14 bound to socket 1[core 14[hwt 0-1]]:
[../..../BB/..][../BB/..../..../..]
[compute-641-8.local:04660] MCW rank 15 bound to socket 1[core 15[hwt 0-1]]:
[../..../BB/..][../BB/..../..../..]
[compute-641-20.local:15368] MCW rank 16 bound to socket 0[core 0[hwt 0-1]]:
[BB/..../..../..][../..../..../..]
[compute-641-20.local:15368] MCW rank 17 bound to socket 0[core 1[hwt 0-1]]:
[../BB/..../..../..][../..../..../..]
[compute-641-20.local:15368] MCW rank 18 bound to socket 0[core 2[hwt 0-1]]:
[../BB/..../..../..][../..../..../..]
[compute-641-20.local:15368] MCW rank 19 bound to socket 0[core 3[hwt 0-1]]:
[../BB/..../..../..][../..../..../..]
[compute-641-20.local:15368] MCW rank 20 bound to socket 0[core 4[hwt 0-1]]:

```



```

[../..../BB/../../][../..../..../..../]
[compute-641-20.local:15368] MCW rank 21 bound to socket 0[core 5[hwt 0-1]]:
[../..../BB/../../][../..../..../..../]
[compute-641-20.local:15368] MCW rank 22 bound to socket 0[core 6[hwt 0-1]]:
[../..../BB/../../][../..../..../..../]
[compute-641-20.local:15368] MCW rank 23 bound to socket 0[core 7[hwt 0-1]]:
[../..../BB/../../][../..../..../..../]
[compute-641-20.local:15368] MCW rank 24 bound to socket 1[core 8[hwt 0-1]]:
[../..../BB/../../][BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 25 bound to socket 1[core 9[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 26 bound to socket 1[core 10[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 27 bound to socket 1[core 11[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 28 bound to socket 1[core 12[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 29 bound to socket 1[core 13[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 30 bound to socket 1[core 14[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]
[compute-641-20.local:15368] MCW rank 31 bound to socket 1[core 15[hwt 0-1]]:
[../..../BB/../../][../BB/../../..../..../]

```

## C Algorithm Efficiency




Figure 6: Algorithm's Efficiency per #process

## D Algorithm Load Balance Data

|     | Core - AllReduce - 1966080 |         |        |        |        |        |        |
|-----|----------------------------|---------|--------|--------|--------|--------|--------|
| P1  | 3250225                    | 1773115 | 973711 | 627109 | 463435 | 440398 | 419407 |
| P2  | 3250228                    | 1772714 | 973413 | 626933 | 462724 | 441614 | 419194 |
| P3  | X                          | 1773034 | 973340 | 626604 | 462799 | 440323 | 419984 |
| P4  | X                          | 1772687 | 973459 | 626168 | 462175 | 442058 | 418862 |
| P5  | X                          | X       | 973083 | 626173 | 462340 | 441660 | 419297 |
| P6  | X                          | X       | 973139 | 626069 | 462194 | 440175 | 418218 |
| P7  | X                          | X       | 972885 | 626059 | 462350 | 441756 | 416986 |
| P8  | X                          | X       | 973190 | 684428 | 465260 | 437604 | 415664 |
| P9  | X                          | X       | X      | 684601 | 463795 | 438268 | 417197 |
| P10 | X                          | X       | X      | 626825 | 462869 | 438323 | 416684 |
| P11 | X                          | X       | X      | 684599 | 463348 | 436440 | 415836 |
| P12 | X                          | X       | X      | 684722 | 463745 | 438373 | 419485 |
| P13 | X                          | X       | X      | X      | 463088 | 438238 | 419174 |
| P14 | X                          | X       | X      | X      | 463060 | 438374 | 419421 |
| P15 | X                          | X       | X      | X      | 463086 | 438125 | 419515 |
| P16 | X                          | X       | X      | X      | 463438 | 451492 | 430961 |
| P17 | X                          | X       | X      | X      | X      | 449151 | 427549 |
| P18 | X                          | X       | X      | X      | X      | 445565 | 425934 |
| P19 | X                          | X       | X      | X      | X      | 444617 | 423470 |
| P20 | X                          | X       | X      | X      | X      | 444137 | 423015 |
| P21 | X                          | X       | X      | X      | X      | 445233 | 423530 |
| P22 | X                          | X       | X      | X      | X      | 445975 | 424788 |
| P23 | X                          | X       | X      | X      | X      | 447211 | 425460 |
| P24 | X                          | X       | X      | X      | X      | 456848 | 439437 |
| P25 | X                          | X       | X      | X      | X      | X      | 434067 |
| P26 | X                          | X       | X      | X      | X      | X      | 429440 |
| P27 | X                          | X       | X      | X      | X      | X      | 428112 |
| P28 | X                          | X       | X      | X      | X      | X      | 428037 |
| P29 | X                          | X       | X      | X      | X      | X      | 428467 |
| P30 | X                          | X       | X      | X      | X      | X      | 429372 |
| P31 | X                          | X       | X      | X      | X      | X      | 431790 |
| P32 | X                          | X       | X      | X      | X      | X      | 442305 |

Figure 7: Load Balance using L3 size data set, Reduce+Broadcast strategy and map by Node , time in microseconds

|     | Core - AllReduce - 62914560 |          |          |          |          |          |          |
|-----|-----------------------------|----------|----------|----------|----------|----------|----------|
| P1  | 161036295                   | 86903623 | 47109528 | 31552725 | 24556621 | 20064670 | 17661518 |
| P2  | 161036272                   | 86903302 | 47110309 | 31553172 | 24555431 | 20067490 | 17654873 |
| P3  | X                           | 86903602 | 47109988 | 31551525 | 24554672 | 20062099 | 17655962 |
| P4  | X                           | 86903415 | 47109397 | 31551204 | 24554053 | 20062325 | 17654505 |
| P5  | X                           | X        | 47109253 | 31551154 | 24554340 | 20068430 | 17663347 |
| P6  | X                           | X        | 47109643 | 31550735 | 24554593 | 20068649 | 17663313 |
| P7  | X                           | X        | 47109673 | 31550827 | 24554901 | 20069067 | 17663703 |
| P8  | X                           | X        | 47109670 | 31614591 | 24558898 | 20068447 | 17655587 |
| P9  | X                           | X        | X        | 31614057 | 24556431 | 20064589 | 17645346 |
| P10 | X                           | X        | X        | 31613559 | 24554585 | 20048186 | 17637659 |
| P11 | X                           | X        | X        | 31613406 | 24554226 | 20046218 | 17649617 |
| P12 | X                           | X        | X        | 31550456 | 24554414 | 20061102 | 17649142 |
| P13 | X                           | X        | X        | X        | 24554743 | 20061019 | 17649595 |
| P14 | X                           | X        | X        | X        | 24554733 | 20061336 | 17649584 |
| P15 | X                           | X        | X        | X        | 24556432 | 20061479 | 17649440 |
| P16 | X                           | X        | X        | X        | 24554742 | 20110732 | 17729604 |
| P17 | X                           | X        | X        | X        | X        | 20105073 | 17731499 |
| P18 | X                           | X        | X        | X        | X        | 20092806 | 17729184 |
| P19 | X                           | X        | X        | X        | X        | 20092534 | 17724946 |
| P20 | X                           | X        | X        | X        | X        | 20092555 | 17739526 |
| P21 | X                           | X        | X        | X        | X        | 20093093 | 17729173 |
| P22 | X                           | X        | X        | X        | X        | 20111303 | 17726664 |
| P23 | X                           | X        | X        | X        | X        | 20110182 | 17725191 |
| P24 | X                           | X        | X        | X        | X        | 20137692 | 17715527 |
| P25 | X                           | X        | X        | X        | X        | X        | 17724934 |
| P26 | X                           | X        | X        | X        | X        | X        | 17724229 |
| P27 | X                           | X        | X        | X        | X        | X        | 17719171 |
| P28 | X                           | X        | X        | X        | X        | X        | 17728725 |
| P29 | X                           | X        | X        | X        | X        | X        | 17724437 |
| P30 | X                           | X        | X        | X        | X        | X        | 17717284 |
| P31 | X                           | X        | X        | X        | X        | X        | 17718655 |
| P32 | X                           | X        | X        | X        | X        | X        | 17731028 |

Figure 8: Load Balance using RAM size data set, Reduce+Broadcast strategy and map by Node, time in microseconds

|     | Core - ReduceBcast - 1966080 |         |        |        |        |        |        |
|-----|------------------------------|---------|--------|--------|--------|--------|--------|
| P1  | 3155165                      | 1773472 | 969258 | 629323 | 463229 | 457348 | 426766 |
| P2  | 3155159                      | 1773300 | 969267 | 629085 | 463575 | 455275 | 425985 |
| P3  | X                            | 1773329 | 968355 | 628792 | 463863 | 455341 | 428254 |
| P4  | X                            | 1773150 | 967825 | 628633 | 463768 | 453822 | 428263 |
| P5  | X                            | X       | 967844 | 628562 | 463448 | 456861 | 427903 |
| P6  | X                            | X       | 967869 | 628682 | 463236 | 456830 | 421403 |
| P7  | X                            | X       | 967419 | 628531 | 463248 | 456727 | 425675 |
| P8  | X                            | X       | 967891 | 629650 | 465361 | 455240 | 427316 |
| P9  | X                            | X       | X      | 685234 | 464185 | 454487 | 424097 |
| P10 | X                            | X       | X      | 685122 | 464263 | 456525 | 424099 |
| P11 | X                            | X       | X      | 685000 | 463539 | 454288 | 422376 |
| P12 | X                            | X       | X      | 685183 | 463298 | 456141 | 423693 |
| P13 | X                            | X       | X      | X      | 463325 | 455947 | 423801 |
| P14 | X                            | X       | X      | X      | 463330 | 456260 | 424189 |
| P15 | X                            | X       | X      | X      | 464449 | 456018 | 424031 |
| P16 | X                            | X       | X      | X      | 463483 | 460912 | 434469 |
| P17 | X                            | X       | X      | X      | X      | 463963 | 432236 |
| P18 | X                            | X       | X      | X      | X      | 461663 | 426899 |
| P19 | X                            | X       | X      | X      | X      | 457791 | 426685 |
| P20 | X                            | X       | X      | X      | X      | 458454 | 427310 |
| P21 | X                            | X       | X      | X      | X      | 459891 | 428651 |
| P22 | X                            | X       | X      | X      | X      | 458462 | 429390 |
| P23 | X                            | X       | X      | X      | X      | 460641 | 431757 |
| P24 | X                            | X       | X      | X      | X      | 471793 | 444969 |
| P25 | X                            | X       | X      | X      | X      | X      | 440127 |
| P26 | X                            | X       | X      | X      | X      | X      | 432119 |
| P27 | X                            | X       | X      | X      | X      | X      | 430071 |
| P28 | X                            | X       | X      | X      | X      | X      | 430724 |
| P29 | X                            | X       | X      | X      | X      | X      | 431578 |
| P30 | X                            | X       | X      | X      | X      | X      | 432446 |
| P31 | X                            | X       | X      | X      | X      | X      | 433717 |
| P32 | X                            | X       | X      | X      | X      | X      | 450662 |

Figure 9: Load Balance using L3 size data set, Allreduce strategy and map by Node, time in microseconds

|     | Core - ReduceBcast - 62914560 |          |          |          |          |          |          |
|-----|-------------------------------|----------|----------|----------|----------|----------|----------|
| P1  | 161043479                     | 86514848 | 47501097 | 31445145 | 24232983 | 19922419 | 17513921 |
| P2  | 161043479                     | 86515047 | 47500784 | 31444821 | 24231747 | 19913484 | 17518168 |
| P3  | X                             | 86514986 | 47500782 | 31444295 | 24230803 | 19923246 | 17521297 |
| P4  | X                             | 86514874 | 47500174 | 31443949 | 24230478 | 19918035 | 17522094 |
| P5  | X                             | X        | 47500072 | 31443558 | 24230583 | 19924570 | 17521869 |
| P6  | X                             | X        | 47500157 | 31444069 | 24230692 | 19923958 | 17522383 |
| P7  | X                             | X        | 47500080 | 31443946 | 24231144 | 19924761 | 17522524 |
| P8  | X                             | X        | 47500116 | 31501773 | 24236718 | 19918781 | 17514292 |
| P9  | X                             | X        | X        | 31501621 | 24233419 | 19898834 | 17506906 |
| P10 | X                             | X        | X        | 31501374 | 24231317 | 19905524 | 17497493 |
| P11 | X                             | X        | X        | 31501286 | 24231905 | 19911501 | 17515602 |
| P12 | X                             | X        | X        | 31444044 | 24231744 | 19911038 | 17515644 |
| P13 | X                             | X        | X        | X        | 24232067 | 19910336 | 17516258 |
| P14 | X                             | X        | X        | X        | 24232127 | 19910098 | 17515932 |
| P15 | X                             | X        | X        | X        | 24232183 | 19910120 | 17515908 |
| P16 | X                             | X        | X        | X        | 24233318 | 20114950 | 17593285 |
| P17 | X                             | X        | X        | X        | X        | 20103959 | 17591275 |
| P18 | X                             | X        | X        | X        | X        | 20101632 | 17590966 |
| P19 | X                             | X        | X        | X        | X        | 20095424 | 17603031 |
| P20 | X                             | X        | X        | X        | X        | 20101151 | 17592016 |
| P21 | X                             | X        | X        | X        | X        | 20100805 | 17573031 |
| P22 | X                             | X        | X        | X        | X        | 20125474 | 17584302 |
| P23 | X                             | X        | X        | X        | X        | 20146462 | 17592072 |
| P24 | X                             | X        | X        | X        | X        | 20122313 | 17600156 |
| P25 | X                             | X        | X        | X        | X        | X        | 17588067 |
| P26 | X                             | X        | X        | X        | X        | X        | 17585896 |
| P27 | X                             | X        | X        | X        | X        | X        | 17588863 |
| P28 | X                             | X        | X        | X        | X        | X        | 17580339 |
| P29 | X                             | X        | X        | X        | X        | X        | 17573667 |
| P30 | X                             | X        | X        | X        | X        | X        | 17577871 |
| P31 | X                             | X        | X        | X        | X        | X        | 17578234 |
| P32 | X                             | X        | X        | X        | X        | X        | 17592524 |

Figure 10: Load Balance using RAM size data set, Allreduce strategy and map by Node, time in microseconds

|     | Node - AllReduce - 1966080 |         |         |        |        |        |        |
|-----|----------------------------|---------|---------|--------|--------|--------|--------|
| P1  | 3231829                    | 1821536 | 1059013 | 801655 | 635935 | 506297 | 414972 |
| P2  | 3231527                    | 1816871 | 1059890 | 801138 | 630146 | 502961 | 415135 |
| P3  | X                          | 1819714 | 1058749 | 801263 | 627646 | 519818 | 414517 |
| P4  | X                          | 1822751 | 1061166 | 801104 | 629996 | 503805 | 408036 |
| P5  | X                          | X       | 1056222 | 805210 | 629514 | 509146 | 410984 |
| P6  | X                          | X       | 1056193 | 801396 | 629894 | 507517 | 409270 |
| P7  | X                          | X       | 1058738 | 801124 | 663297 | 500848 | 406187 |
| P8  | X                          | X       | 1070490 | 799820 | 633851 | 501307 | 405575 |
| P9  | X                          | X       | X       | 798066 | 640142 | 502101 | 411021 |
| P10 | X                          | X       | X       | 803408 | 634748 | 502972 | 404724 |
| P11 | X                          | X       | X       | 804454 | 636100 | 503457 | 411568 |
| P12 | X                          | X       | X       | 816201 | 636252 | 555165 | 406693 |
| P13 | X                          | X       | X       | X      | 636537 | 503835 | 413451 |
| P14 | X                          | X       | X       | X      | 639326 | 503365 | 408818 |
| P15 | X                          | X       | X       | X      | 641984 | 552707 | 421303 |
| P16 | X                          | X       | X       | X      | 646885 | 550597 | 409550 |
| P17 | X                          | X       | X       | X      | X      | 504925 | 417559 |
| P18 | X                          | X       | X       | X      | X      | 507141 | 420760 |
| P19 | X                          | X       | X       | X      | X      | 554562 | 410092 |
| P20 | X                          | X       | X       | X      | X      | 556582 | 412163 |
| P21 | X                          | X       | X       | X      | X      | 555876 | 409823 |
| P22 | X                          | X       | X       | X      | X      | 555418 | 410641 |
| P23 | X                          | X       | X       | X      | X      | 554854 | 409554 |
| P24 | X                          | X       | X       | X      | X      | 517967 | 408006 |
| P25 | X                          | X       | X       | X      | X      | X      | 410303 |
| P26 | X                          | X       | X       | X      | X      | X      | 407059 |
| P27 | X                          | X       | X       | X      | X      | X      | 411123 |
| P28 | X                          | X       | X       | X      | X      | X      | 412023 |
| P29 | X                          | X       | X       | X      | X      | X      | 407639 |
| P30 | X                          | X       | X       | X      | X      | X      | 408557 |
| P31 | X                          | X       | X       | X      | X      | X      | 410155 |
| P32 | X                          | X       | X       | X      | X      | X      | 428619 |

Figure 11: Load Balance using L3 size data set, Reduce + Broadcast strategy and map by Core, time in microseconds

|     | Node - AllReduce - 62914560 |          |          |          |          |          |          |
|-----|-----------------------------|----------|----------|----------|----------|----------|----------|
| P1  | 165669824                   | 88673686 | 50697079 | 35301892 | 30006547 | 21629653 | 17669833 |
| P2  | 165668525                   | 88669452 | 50694332 | 35306792 | 29985782 | 21614941 | 17678305 |
| P3  | X                           | 88672332 | 50696845 | 35308445 | 29978377 | 21619709 | 17676724 |
| P4  | X                           | 88674666 | 50690407 | 35309697 | 29972569 | 21606586 | 17670924 |
| P5  | X                           | X        | 50689378 | 35311755 | 29977817 | 21616073 | 17674100 |
| P6  | X                           | X        | 50685830 | 35990818 | 29975799 | 21617664 | 17673696 |
| P7  | X                           | X        | 50688896 | 35995895 | 29976594 | 21624602 | 17674380 |
| P8  | X                           | X        | 50711100 | 36331200 | 30035776 | 21674330 | 17739876 |
| P9  | X                           | X        | X        | 36336452 | 30022202 | 21665539 | 17741001 |
| P10 | X                           | X        | X        | 36343867 | 30021443 | 21665492 | 17735519 |
| P11 | X                           | X        | X        | 36350064 | 30029102 | 21679318 | 17719898 |
| P12 | X                           | X        | X        | 36349503 | 30017281 | 21685348 | 17719166 |
| P13 | X                           | X        | X        | X        | 30054683 | 21684497 | 17717961 |
| P14 | X                           | X        | X        | X        | 30026517 | 21683520 | 17718203 |
| P15 | X                           | X        | X        | X        | 30027522 | 21680122 | 17727836 |
| P16 | X                           | X        | X        | X        | 30044368 | 21681035 | 17746496 |
| P17 | X                           | X        | X        | X        | X        | 21684465 | 17729964 |
| P18 | X                           | X        | X        | X        | X        | 21663192 | 17749670 |
| P19 | X                           | X        | X        | X        | X        | 21686930 | 17744322 |
| P20 | X                           | X        | X        | X        | X        | 21678638 | 17750931 |
| P21 | X                           | X        | X        | X        | X        | 21677458 | 17751519 |
| P22 | X                           | X        | X        | X        | X        | 21672803 | 17727006 |
| P23 | X                           | X        | X        | X        | X        | 21672027 | 17728112 |
| P24 | X                           | X        | X        | X        | X        | 21682548 | 17745140 |
| P25 | X                           | X        | X        | X        | X        | X        | 17742407 |
| P26 | X                           | X        | X        | X        | X        | X        | 17715518 |
| P27 | X                           | X        | X        | X        | X        | X        | 17742753 |
| P28 | X                           | X        | X        | X        | X        | X        | 17737345 |
| P29 | X                           | X        | X        | X        | X        | X        | 17743104 |
| P30 | X                           | X        | X        | X        | X        | X        | 17729228 |
| P31 | X                           | X        | X        | X        | X        | X        | 17731266 |
| P32 | X                           | X        | X        | X        | X        | X        | 17745877 |

Figure 12: Load Balance using RAM size data set, Reduce + Broadcast strategy and map by Core, time in microseconds



|     | Node - ReduceBcast - 1966080 |         |         |         |         |         |         |
|-----|------------------------------|---------|---------|---------|---------|---------|---------|
| P1  | 4769571                      | 3319116 | 2852045 | 2795192 | 2626630 | 2590025 | 3486717 |
| P2  | 4763137                      | 3218843 | 2855613 | 2574420 | 2628842 | 2579839 | 3482236 |
| P3  | X                            | 3398992 | 2665578 | 2575431 | 2632587 | 2567779 | 3484764 |
| P4  | X                            | 3399891 | 2486828 | 2614688 | 2591595 | 2559504 | 3486592 |
| P5  | X                            | X       | 3096521 | 2527876 | 2591737 | 2561690 | 3482755 |
| P6  | X                            | X       | 3078220 | 2870354 | 2586427 | 2561645 | 3484954 |
| P7  | X                            | X       | 2874649 | 2603084 | 2588170 | 2566731 | 3930841 |
| P8  | X                            | X       | 2875481 | 2861190 | 2453687 | 2596578 | 3931305 |
| P9  | X                            | X       | X       | 2685506 | 2381884 | 2145319 | 3687572 |
| P10 | X                            | X       | X       | 2660669 | 2607978 | 2144296 | 3685494 |
| P11 | X                            | X       | X       | 2571919 | 2360296 | 2111990 | 3688690 |
| P12 | X                            | X       | X       | 2814330 | 2909931 | 2155428 | 3699763 |
| P13 | X                            | X       | X       | X       | 2617342 | 2115365 | 3239476 |
| P14 | X                            | X       | X       | X       | 2835887 | 2605731 | 3241517 |
| P15 | X                            | X       | X       | X       | 2439940 | 2547792 | 3239300 |
| P16 | X                            | X       | X       | X       | 2710503 | 2806348 | 2802772 |
| P17 | X                            | X       | X       | X       | X       | 2615661 | 2803609 |
| P18 | X                            | X       | X       | X       | X       | 2820762 | 2561643 |
| P19 | X                            | X       | X       | X       | X       | 2851583 | 2563217 |
| P20 | X                            | X       | X       | X       | X       | 2165686 | 3683784 |
| P21 | X                            | X       | X       | X       | X       | 2304310 | 3678261 |
| P22 | X                            | X       | X       | X       | X       | 2628671 | 3681018 |
| P23 | X                            | X       | X       | X       | X       | 2541785 | 3683220 |
| P24 | X                            | X       | X       | X       | X       | 2575586 | 3239388 |
| P25 | X                            | X       | X       | X       | X       | X       | 3240542 |
| P26 | X                            | X       | X       | X       | X       | X       | 3239705 |
| P27 | X                            | X       | X       | X       | X       | X       | 2802709 |
| P28 | X                            | X       | X       | X       | X       | X       | 2803605 |
| P29 | X                            | X       | X       | X       | X       | X       | 2554067 |
| P30 | X                            | X       | X       | X       | X       | X       | 2554844 |
| P31 | X                            | X       | X       | X       | X       | X       | 2556066 |
| P32 | X                            | X       | X       | X       | X       | X       | 2765386 |

Figure 13: Load Balance using L3 size data set, Allreduce strategy and map by Core, time in microseconds

|     | Node - ReduceBcast - 62914560 |          |          |          |          |          |          |
|-----|-------------------------------|----------|----------|----------|----------|----------|----------|
| P1  | 166102740                     | 88960723 | 49696011 | 36219776 | 29716549 | 20897737 | 17673362 |
| P2  | 166103715                     | 88958579 | 49697891 | 36220140 | 29701424 | 20900162 | 17660936 |
| P3  | X                             | 88960066 | 49699290 | 36219901 | 29699057 | 20903363 | 17670917 |
| P4  | X                             | 88964274 | 50741800 | 36217366 | 29702582 | 20909084 | 17673707 |
| P5  | X                             | X        | 50733685 | 36220259 | 29705891 | 20911422 | 17672952 |
| P6  | X                             | X        | 50735062 | 36258872 | 29701023 | 20912716 | 17673383 |
| P7  | X                             | X        | 50737021 | 36253880 | 29704897 | 20935616 | 17675606 |
| P8  | X                             | X        | 50738769 | 36256008 | 29770103 | 20982735 | 17738255 |
| P9  | X                             | X        | X        | 36241614 | 29770668 | 20986214 | 17733260 |
| P10 | X                             | X        | X        | 36241149 | 29751542 | 21019368 | 17739864 |
| P11 | X                             | X        | X        | 36258416 | 29754914 | 21071065 | 17739348 |
| P12 | X                             | X        | X        | 36243437 | 29772451 | 21408312 | 17756759 |
| P13 | X                             | X        | X        | X        | 29751555 | 21417829 | 17761854 |
| P14 | X                             | X        | X        | X        | 29748729 | 21636141 | 17771440 |
| P15 | X                             | X        | X        | X        | 29754814 | 21712246 | 17771898 |
| P16 | X                             | X        | X        | X        | 29766501 | 21758064 | 17708460 |
| P17 | X                             | X        | X        | X        | X        | 21770251 | 17690599 |
| P18 | X                             | X        | X        | X        | X        | 21817116 | 17782136 |
| P19 | X                             | X        | X        | X        | X        | 21824528 | 17730952 |
| P20 | X                             | X        | X        | X        | X        | 21904103 | 17732214 |
| P21 | X                             | X        | X        | X        | X        | 21925369 | 17739886 |
| P22 | X                             | X        | X        | X        | X        | 21988943 | 17728491 |
| P23 | X                             | X        | X        | X        | X        | 21998948 | 17740739 |
| P24 | X                             | X        | X        | X        | X        | 21961429 | 17717975 |
| P25 | X                             | X        | X        | X        | X        | X        | 17848321 |
| P26 | X                             | X        | X        | X        | X        | X        | 17723676 |
| P27 | X                             | X        | X        | X        | X        | X        | 17791483 |
| P28 | X                             | X        | X        | X        | X        | X        | 17718393 |
| P29 | X                             | X        | X        | X        | X        | X        | 17743480 |
| P30 | X                             | X        | X        | X        | X        | X        | 17727001 |
| P31 | X                             | X        | X        | X        | X        | X        | 17726600 |
| P32 | X                             | X        | X        | X        | X        | X        | 17740825 |

Figure 14: Load Balance using RAM size data set, Allreduce strategy and map by Core, time in microseconds